

Future-Proofing TSA Operations: Data-Driven Resource Optimization

```
library(readxl)
#install.packages('fpp3')
#install.packages("urca")
library(fpp3)

Registered S3 method overwritten by 'tsibble':
  method           from
  as_tibble.grouped_df dplyr

-- Attaching packages ----- fpp3 1.0.2 --

v tibble      3.2.1    v tsibble     1.1.6
v dplyr       1.1.4    v tsibbledata 0.4.1
v tidyr       1.3.1    v feasts       0.4.2
v lubridate   1.9.4    v fable        0.4.1
v ggplot2     3.5.1

-- Conflicts ----- fpp3_conflicts --
x lubridate::date()    masks base::date()
x dplyr::filter()      masks stats::filter()
x tsibble::intersect() masks base::intersect()
x tsibble::interval()  masks lubridate::interval()
x dplyr::lag()         masks stats::lag()
x tsibble::setdiff()   masks base::setdiff()
x tsibble::union()     masks base::union()

library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
vforcats 1.0.0      vreadr    2.1.5
vpurrr    1.0.2      vstringr 1.5.1

-- Conflicts ----- tidyverse_conflicts() --
xdplyr::filter()    masks stats::filter()
xtsibble::interval() masks lubridate::interval()
xdplyr::lag()        masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become non-conflicting
```

```
library(fpp3)
library(slider)
library(urca)
library(scales)
```

Attaching package: 'scales'

The following object is masked from 'package:purrr':

discard

The following object is masked from 'package:readr':

col_factor

Exploratory Data Analysis


```
tsa_clean <- read_excel(here::here("Data", "TSA.xlsx"))
```

```
head(tsa_clean)
```

```
# A tibble: 6 x 2
  Date                Numbers
  <dttm>              <dbl>
1 2025-11-06 00:00:00 2703787
2 2025-11-05 00:00:00 2162183
3 2025-11-04 00:00:00 2015297
4 2025-11-03 00:00:00 2492532
```

```
5 2025-11-02 00:00:00 2279000
6 2025-11-01 00:00:00 2029517
```

```
library(ggplot2)
autoplot(tsa_clean |> mutate(Date = as.Date(Date)) |> as_tsibble(index = Date) , Numbers) +
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale())) +
  labs(
    title = "TSA Daily Passenger Volume",
    y = "Passengers",
    x = NULL
  ) +
  theme_minimal()
```



```
# Glimpse at dataset
glimpse(tsa_clean)
```

```
Rows: 2,502
Columns: 2
$ Date      <dttm> 2025-11-06, 2025-11-05, 2025-11-04, 2025-11-03, 2025-11-02, 2-
$ Numbers   <dbl> 2703787, 2162183, 2015297, 2492532, 2279000, 2029517, 2678000, ~
```

```
# Number Column summary  
summary(tsa_clean$Numbers)
```

```
Min. 1st Qu. Median Mean 3rd Qu. Max.  
113147 1811221 2188235 2024965 2498146 3096797
```

```
# Missing data check  
sum(is.na(tsa_clean$Numbers))
```

```
[1] 0
```


```
# Range of the Dates  
min(tsa_clean$Date); max(tsa_clean$Date)
```

```
[1] "2019-01-01 UTC"
```

```
[1] "2025-11-06 UTC"
```

```
# Outlier analysis  
boxplot(tsa_clean$Numbers, main = "Outlier Analysis", ylab = "Values", horizontal = TRUE)
```

Outlier Analysis



```

# List of outliers -- 2020-2021
Q1 <- quantile(tsa_clean$Numbers, 0.25, na.rm = TRUE)
Q3 <- quantile(tsa_clean$Numbers, 0.75, na.rm = TRUE)
IQR_value <- IQR(tsa_clean$Numbers, na.rm = TRUE)

lower_bound <- Q1 - (1.5 * IQR_value)
upper_bound <- Q3 + (1.5 * IQR_value)

outliers <- tsa_clean |> filter(Numbers < lower_bound | Numbers > upper_bound)

head(outliers)

```

```

# A tibble: 6 x 2
  Date           Numbers
  <dttm>        <dbl>
1 2021-01-06 00:00:00 705249
2 2021-01-09 00:00:00 750419
3 2021-01-11 00:00:00 750407
4 2021-01-12 00:00:00 557517
5 2021-01-13 00:00:00 605887
6 2021-01-16 00:00:00 729703


```

```

tsa_monthly <- tsa_clean |>
  mutate(Month = yearmonth(Date)) |>
  group_by(Month) |>
  summarise(Total = sum(Numbers, na.rm = TRUE)) |>
  ungroup() |>
  as_tsibble(index = Month)

autoplot(tsa_monthly, Total) +
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale())) +
  labs(
    title = "TSA Monthly Passenger Volume",
    y = "Passengers",
    x = NULL
  ) +
  theme_minimal()

```




Training/Test

Convert to TSIBBLE object, filter the months, and create a Totals Column.

```
covid_window <- tsa_monthly %>% # Changed from tsa_clean to tsa_monthly
  filter_index("2019 Dec" ~ "2021 Dec")

# Plot the Drop
covid_window %>%
  autoplot(Total, linewidth = 1) +
  # Formatting axes
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale())) +
  scale_x_yearmonth(date_breaks = "1 month", date_labels = "%b %Y") +
  labs(
    title = "TSA Passenger Volume: The COVID-19 Shock",
    subtitle = "Analysis of Drop from Dec 2019 to Dec 2021",
    y = "Monthly Passengers",
    x = NULL
  ) +
```


```
theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



```
# Expect columns: Date, Numbers (daily)
tsa_clean <- tsa_monthly %>%
  filter(Month < max(Month)) %>%
  filter(year(Month) >= 2019) %>%
  mutate(
    covid_shock = if_else(Month >= yearmonth("2020 Feb") & Month <= yearmonth("2021 Dec"), 1,
  )
  )

autoplot(tsa_clean, Total) + labs(title = "TSA monthly totals", x = NULL, y = "Total")+
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale()))
```

TSA monthly totals




```
head(tsa_clean)
```

```
# A tsibble: 6 x 3 [1M]
  Month      Total covid_shock
  <mth>    <dbl>     <dbl>
1 2019 Jan  61694899      0
2 2019 Feb  58535547      0
3 2019 Mar  72714771      0
4 2019 Apr  69754997      0
5 2019 May  74582398      0
6 2019 Jun  76489355      0
```

Monthly Timeseries Plot

```
tsa_clean %>%
  autoplot(Total) +
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale()))+
  labs(
  title = "TSA Monthly Passenger Volume",
  y = "Total Passengers"
```

```
) +  
theme_minimal()
```




STL Decomposition

```
tsa_clean %>%  
  model(STL(Total)) %>%  
  components() %>%  
  autoplot() +  
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale()))
```

STL decomposition

Total = trend + season_year + remainder



Training/Testing Phase

```
h <- 12 # goo for evalaution seasonality

# train test, fit
n_total <- nrow(tsa_clean)
train <- tsa_clean %>% filter(year(Month) <= 2023)
test  <- tsa_clean %>% filter(year(Month) > 2023)
```

```
max(test$Month) ; min(test$Month)
```

```
<yearmonth[1]>
[1] "2025 Oct"
```

```
<yearmonth[1]>
[1] "2024 Jan"
```

```
max(train$Month) ; min(train$Month)
```

```

<yearmonth[1]>
[1] "2023 Dec"

<yearmonth[1]>
[1] "2019 Jan"


# Simple diagnostic models (not your final comparison yet)
diag_fit <- tsa_clean %>%
  model(
    ETS_raw    = ETS(Total),
    ETS_log    = ETS(log(Total)),
    ARIMA_raw  = ARIMA(Total),
    ARIMA_log  = ARIMA(log(Total))
  )

# Residuals vs fitted for both
diag_resids <- diag_fit %>%
  augment()

diag_resids %>%
  ggplot(aes(x = .fitted, y = .resid)) +
  geom_point(alpha = 0.5) +
  geom_hline(yintercept = 0, linetype = "dashed") +
  facet_wrap(~ .model, scales = "free_x") +
  labs(
    title = "Residuals vs Fitted: ETS/ARIMA on Raw vs Log(Total)",
    x = "Fitted values",
    y = "Residuals"
  ) +
  scale_x_continuous(
    labels = scales::label_number(scale_cut = scales::cut_short_scale())
  ) +
  theme_minimal()

```

Residuals vs Fitted: ETS/ARIMA on Raw vs Log(Total)



Residuals vs fitted plots for ETS and ARIMA on both the original and log-transformed scales show that the raw-scale models (ETS(Total) and ARIMA(Total)) have tighter, more homoscedastic residuals centered around zero. In contrast, the log-transformed models exhibit larger and more structured residuals. Together with the poorer RMSE/MAE for the log models, this indicates that a log transformation does not improve the fit and is not appropriate for this series. We therefore proceed using models on the original scale only.

Fitting Stage

```

fit <- train %>%
  model(
    ETS    = ETS(Total),
    ARIMA = ARIMA(Total),
    ARIMA_pd = ARIMA((Total) ~ covid_shock + season() + trend())
  )

fc <- new_data(train, n = nrow(test)) %>%
  mutate(covid_shock = 0)


fc <- fit %>%
  forecast(new_data = fc)

```

```

fc %>%
  filter(year(Month) >= 2023) %>%
  autoplot(tsa_clean %>% filter(year(Month) >= 2023), level = NULL) +
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale())) +
  labs(
    title = "Forecast Comparison (2023–2026)",
    y = "Passengers",
    x = NULL
  ) +
  theme_minimal()

```



```

acc_baseline_models <- accuracy(fc, test) %>%
  select(.model:MAPE)

acc_baseline_models %>%
  arrange(RMSE) %>%
  knitr::kable()

```

Warning in attr(x, "align"): 'xfun::attr()' is deprecated.
 Use 'xfun::attr2()' instead.
 See help("Deprecated")

```

Warning in attr(x, "format"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

```

.model	.type	ME	RMSE	MAE	MPE	MAPE
ETS	Test	2939451	7018582	6124212	3.160949	8.088764
ARIMA_pd	Test	7160362	8266266	7223109	9.146159	9.243218
ARIMA	Test	12958056	15346419	14031963	16.394367	18.057124

ARIMA Set-Up

```
print('Stationary Test')
```

```
[1] "Stationary Test"
```

```
train |> fabletools::features(Total, c(unitroot_kpss, unitroot_ndiffs, unitroot_nsdiffs)) |>
```

```

Warning in attr(x, "align"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

```

```

Warning in attr(x, "format"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

```

kpss_stat	kpss_pvalue	ndiffs	nsdiffs
0.3902161	0.0813724	0	0

Observation: Null hypothesis cannot be rejected at 0.05. Therefore, series is stationary. 1 regular and 1 seasonal differentiation is required.

```
print('Stationary test after applying a season and regular differentiation')
```

```
[1] "Stationary test after applying a season and regular differentiation"
```

```

train |> fabletools::features(Total |> difference(lag=1) , c(unitroot_kpss, unitroot_ndiffs

Warning in attr(x, "align"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

Warning in attr(x, "format"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

```

kpss_stat	kpss_pvalue	ndiffs	nsdiffs
0.1058268	0.1	0	0

```

#orders1 <- tsa_clean |>
#  mutate(Total_d = difference(difference(Total, lag = 12))) |>
#  features(Total_d, unitroot_kpss)
#print(orders1)

```

Differenced series and ACF/PACF plots

```


# plot differenced series and ACF/PACF plots with non-seasonal and seasonal differencing
train |>
  feasts::gg_tsdisplay(
    Total |> difference(lag = 1),
    plot_type = "partial"
  ) +
  ggplot2::scale_y_continuous(labels = label_number(scale_cut = cut_short_scale()))

```

Warning: `gg_tsdisplay()` was deprecated in feasts 0.4.2.
 i Please use `ggttime::gg_tsdisplay()` instead.

Warning: Removed 1 row containing missing values or values outside the scale range
 (`geom_line()`).

Warning: Removed 1 row containing missing values or values outside the scale range
 (`geom_point()`).



```
# model fit
fitted_manual_arima <- train |>
  model(
    ma1_manual = ARIMA(Total ~ pdq(0, 1, 1) + PDQ(0, 1, 1)),
    ar1_manual = ARIMA(Total ~ pdq(1, 1, 0) + PDQ(1, 1, 0))
  )

h_test <- nrow(test)

# Holds both forecast models
fc_manual_arima <- fitted_manual_arima |>
  forecast(h = h_test)

# Manual ARIMA scores from Training Data
fitted_manual_arima |> glance() |> select(.model, AICc) |> arrange(AICc)
```

```
# A tibble: 2 x 2
  .model      AICc
  <chr>      <dbl>
1 ma1_manual 1626.
2 ar1_manual 1632.
```

```

all_models <- train |>
  model(
    ETS      = ETS(Total),
    auto_arima = ARIMA(Total),
    ma1_manual = ARIMA(Total ~ pdq(0, 1, 1) + PDQ(0, 1, 1)),
    ar1_manual = ARIMA(Total ~ pdq(1, 1, 0) + PDQ(1, 1, 0))
  )

all_models |> glance() |> select(.model, AICc) |> arrange(AICc)

```

```

# A tibble: 4 x 2
  .model     AICc
  <chr>     <dbl>
1 ma1_manual 1626.
2 ar1_manual 1632.
3 auto_arima 2080.
4 ETS        2158.

```

All performance metrics combined

```

forecast_all_models <- all_models |>
  forecast(h = nrow(test))

accuracy_combined <- forecast_all_models |>
  accuracy(tsa_clean) |>
  select(.model, RMSE, MAE, MAPE) |>
  arrange(RMSE)
accuracy_combined |> knitr::kable()

```

```

Warning in attr(x, "align"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

```


```

Warning in attr(x, "format"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

```

.model	RMSE	MAE	MAPE
ar1_manual	5734254	4719541	6.348599
ETS	7018582	6124212	8.088764
ma1_manual	7687646	6627659	8.481189
auto_arima	15346419	14031963	18.057124

```
# Plot forecasts from all models, restricted to 2023 onward
forecast_all_models %>%
  filter(year(Month) >= 2023) %>%
  autoplot(tsa_clean %>% filter(year(Month) >= 2023), level = NULL) +
  scale_y_continuous(labels = label_number(scale_cut = cut_short_scale())) +
  labs(
    title = "Forecast Comparison (2023–2026)",
    y = "Passengers",
    x = NULL
  ) +
  theme_minimal()
```



Based on test-set RMSE, the ARIMA(1,1,0)(1,1,0) model (ar1_manual) achieved the best out-of-sample performance (RMSE 5.7M), outperforming both ETS and the alternative ARIMA specifications. Although ma1_manual had slightly better AICc on the training data, ar1_manual provided superior accuracy on the holdout period, so we select ar1_manual as

our final model. Log-transformed and pandemic-adjusted models all produced substantially larger RMSE values and were discarded.

Accuracy Comparison with Cross-validation

```

n_total <- nrow(tsa_clean)
train <- tsa_clean |> slice_head(n = n_total - h)
test <- tsa_clean |> slice_tail(n = h)

cv_train <- tsa_clean |> slice_head(n = n_total - h) |> stretch_tsibble(.init=48, .step = 12)

# new cv fits
cv_fits <- cv_train |>
  model(
    ETS      = ETS(Total),
    ARIMA    = ARIMA(Total),
    ma1_manual = ARIMA(Total ~ pdq(0,1,1) + PDQ(0,1,0)),
    ar1_manual = ARIMA(Total ~ pdq(1,1,0) + PDQ(0,0,1))
  )

fc_cv_models <- cv_fits |> forecast(h = "12 months")

accuracy <- accuracy(fc_cv_models, tsa_clean) |> select(.model, .type, ME, RMSE, MAE, MAPE)
accuracy |> knitr::kable()

```

```

Warning in attr(x, "align"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

```

```

Warning in attr(x, "format"): 'xfun::attr()' is deprecated.
Use 'xfun::attr2()' instead.
See help("Deprecated")

```

.model	.type	ME	RMSE	MAE	MAPE
ma1_manual	Test	-1145959	2985843	2177955	3.054853
ar1_manual	Test	3354766	6594746	5904169	8.039470

.model	.type	ME	RMSE	MAE	MAPE
ETS	Test	4095700	7399181	6646001	9.014394
ARIMA	Test	11219474	13945564	13005485	17.332255

Although ETS produced the best accuracy on the single 12-month hold-out set, its performance dropped considerably under rolling-origin cross-validation. This difference occurs because the simple train/test split evaluates the model on only one recent period (a relatively stable year), where ETS performs well. In contrast, time-series cross-validation tests the model across multiple historical forecasting scenarios, including volatile periods such as the COVID-19 shock and the recovery phase. In these more challenging conditions, ETS is less stable, while the manual ARIMA(0,1,1)(0,1,0) model remains consistently accurate. Therefore, cross-validation indicates that the ARIMA model is more robust overall and is the preferred final model.

Workforce Prediction

```
# Workforce Assumptions (Based on TSA lane throughput)
current_agents <- 47500

# Lane capacity per hour (assumed)
pax_per_lane_hour <- 180

# Daily Operating hours per lane
hours_per_day_lane <- 16

# Days per month
days_per_month <- 30

# TSOs required to operate one lane
tsos_per_lane <- 8

# Monthly lane capacity
pax_per_lane_month <- pax_per_lane_hour * hours_per_day_lane * days_per_month

# Share of total workforce assigned to screen (assumption)
screening_duty_fraction <- 0.5

# Active checkpoint workforce
active_screening_tsos <- current_agents * screening_duty_fraction

# Passengers each TSO effectively supports/month
passengers_per_agent_month <- pax_per_lane_month / tsos_per_lane
```

```

# ma1_manual is your chosen final model
fc_final <- all_models |>
  select(ma1_manual) |>
  forecast(h = nrow(test))

buffer_factor <- 1.10 # 10% safety buffer

tsa_agents_forecast <- fc_final |>
  mutate(
    pax_forecast_mean = as.numeric(.mean),
    pax_forecast_high = pax_forecast_mean * buffer_factor,
    agents_required_exact = pax_forecast_high / passengers_per_agent_month,
    agents_required_ceiling = ceiling(agents_required_exact),
    pct_of_screening_staff = agents_required_ceiling / active_screening_tsos
  )

tsa_dual <- tsa_agents_forecast |>
  as_tibble() |>
  transmute(
    Month,
    `Passengers (Forecast + 10% Safety Buffer)` = pax_forecast_high,
    `Agents Required (rounded)` = agents_required_ceiling,
    `% of Screening Workforce Needed` = pct_of_screening_staff
  ) |>
  pivot_longer(
    cols = -Month,
    names_to = "metric",
    values_to = "value"
  ) |>
  drop_na(value)

ggplot(tsa_dual, aes(x = Month, y = value)) +
  geom_line() +
  geom_point() +
  facet_wrap(~ metric, scales = "free_y", ncol = 1) +
  labs(
    title = "TSA Monthly Passenger Forecast and Workforce Demand\nModel: ARIMA(0,1,1)(0,1,0)",
    x = "Month",
    y = NULL
  ) +
  scale_y_continuous(
    labels = scales::label_number(accuracy = 0.01) # <-- removed scale_cut


```

```

) +
theme_minimal() +
theme(
  strip.text = element_text(size = 12, face = "bold")
)

```

TSA Monthly Passenger Forecast and Workforce Demand Model: ARIMA(0,1,1)(0,1,0) with 10% Safety Buffer



Insight: Nearly 200 TSA Officers are paid by the government but work full-time on union matters. These people do not retain certification to perform screening functions. Additionally, in a recent TSA employee survey, over 60% said poor performers are allowed to stay employed and, not surprisingly, continue to not perform.

Source: <https://www.dhs.gov/news/2025/03/07/dhs-ends-collective-bargaining-tsas-transportation-security-officers-enhancing>

Recommendation We recommend that TSA adopt a Monthly Forecasting Cadence anchored by the ARIMA(0,1,1)(0,1,0) model to guide checkpoint workforce planning. This forecasting approach provides TSA with a reliable month-ahead view of expected passenger volumes, enabling proactive and data-driven staffing decisions.

Using TSA's own operational throughput standards (180 passengers per lane per hour, ~8 TSOs required per lane across shifts) and assuming that approximately 50% of the agency's 47,500 officers rotate through screening duties, our workforce model shows that TSA will need only 7,000–8,500 officers per month on screening duty throughout 2024. Even after applying

a 10% safety buffer to account for forecast uncertainty, this represents only 25–36% of the available screening workforce.

This reveals a significant and consistent capacity margin that TSA can strategically leverage. Monthly forecasting enables TSA to:

Adjust screening deployments by month, rather than relying on static annual templates

Reallocate 15–25% of screeners during low-volume periods to training, leave, administrative roles, or airport support

Reduce unnecessary overstaffing during shoulder travel months

Prepare earlier for predictable surges, decreasing reliance on overtime and last-minute staffing during holiday and summer peaks

Enhance TSA well-being by smoothing workload intensity and improving shift predictability

By aligning staffing to model-based demand—and doing so monthly—TSA can maintain its <30-minute wait-time service level while operating more efficiently, reducing operational strain, and improving workforce readiness without reducing headcount.

```
library(dplyr)
library(lubridate)
library(tsibble)
library(scales)
library(knitr)

# Prep data
tsa_yearly <- tsa_clean %>%
  as_tibble() %>%
  ungroup() %>%
  mutate(Year = year(Month)) %>%
  group_by(Year) %>%
  summarise(
    yearly_total = sum(Total, na.rm = TRUE),
    mean_monthly = mean(Total, na.rm = TRUE),
    median_monthly = median(Total, na.rm = TRUE),
    .groups = "drop"
  ) %>%
  arrange(Year) %>%
  mutate(
    yoy_growth = (yearly_total / lag(yearly_total) - 1),
    Year = as.character(Year)
  )
```

```

tsa_overall <- tsa_clean %>%
  as_tibble() %>%
  ungroup() %>%
  summarise(
    Year          = "All years",
    yearly_total = sum(Total, na.rm = TRUE),
    mean_monthly = mean(Total, na.rm = TRUE),
    median_monthly = median(Total, na.rm = TRUE),
    yoy_growth   = NA_real_
  )

tsa_raw_data <- bind_rows(tsa_overall, tsa_yearly)

# Format and display table
tsa_raw_data %>%
  mutate(
    `Total Passengers` = comma(yearly_total), # Adds commas: 1,000,000
    `Avg Monthly`     = comma(mean_monthly, accuracy = 1),
    `Median Monthly`  = comma(median_monthly, accuracy = 1),
    `YoY Growth`      = if_else(is.na(yoy_growth), "-", percent(yoy_growth, accuracy = 0.1))
  ) %>%
  select(Year, `Total Passengers`, `YoY Growth`, `Avg Monthly`, `Median Monthly`) %>%
  kable(
    align = "r",
    caption = "TSA Passenger Throughput (2019–2025)"
  )

```

Warning in attr(x, "align"): 'xfun::attr()' is deprecated.
 Use 'xfun::attr2()' instead.
 See help("Deprecated")

Warning in attr(x, "format"): 'xfun::attr()' is deprecated.
 Use 'xfun::attr2()' instead.
 See help("Deprecated")

Table 6: TSA Passenger Throughput (2019-2025)

Year	Total Passengers	YoY Growth	Avg Monthly	Median Monthly
All years	5,052,780,460	—	61,619,274	67,616,844
2019	848,102,043	—	70,675,170	72,367,869

Year	Total Passengers	YoY Growth	Avg Monthly	Median Monthly
2020	339,774,756	-59.9%	28,314,563	24,618,172
2021	585,250,987	72.2%	48,770,916	54,258,962
2022	760,071,362	29.9%	63,339,280	65,469,214
2023	858,548,196	13.0%	71,545,683	72,240,801
2024	904,068,577	5.3%	75,339,048	76,847,030
2025	756,964,539	-16.3%	75,696,454	78,205,370

```


tsa_clean %>%
  as_tibble() %>%
  mutate(
    Month_Label = month(Month, label = TRUE, abbr = TRUE),
    Month_Num = month(Month),
    Year_Factor = as.factor(year(Month))
  ) %>%
  ggplot(aes(x = Month_Label, y = Total, group = Year_Factor, color = Year_Factor)) +
  geom_line(size = 1) +
  geom_point(size = 2) +
  
  # Highlight 2019 (Baseline) and 2025 (Current)
  scale_color_viridis_d(option = "turbo", direction = -1) +
  
  # Format chart
  scale_y_continuous(labels = label_number(scale = 1e-6, suffix = " M")) +
  
  theme_minimal() +
  labs(
    title = "Seasonal Travel Patterns: Year-over-Year Comparison",
    subtitle = "Comparing monthly volumes across different years",
    x = "",
    y = "Passengers",
    color = "Year"
  )

```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
 i Please use `linewidth` instead.

Seasonal Travel Patterns: Year-over-Year Comparison

Comparing monthly volumes across different years



```

# Max passengers - 85M pax == 8000 agents
pax_per_agent_ratio <- 10625
total_staff_available <- 23530

# Refit ARIMA(0,1,1)(0,1,0) to full dataset
fit_final <- tsa_clean %>%
  model(
    ma1_manual = ARIMA(Total ~ pdq(0,1,1) + PDQ(0,1,0))
  )

# Forecast forward (covering remaining 2025 and all of 2026)
fc_2026 <- fit_final %>%
  forecast(h = "14 months")

# Calculate the Metrics
tsa_agents_forecast <- fc_2026 %>%
  as_tibble() %>%
  mutate(
    pax_forecast_high = .mean * 1.10,
    agents_required_ceiling = ceiling(pax_forecast_high / pax_per_agent_ratio),
    pct_of_screening_staff = agents_required_ceiling / total_staff_available
  ) %>%
  select(Month, pax_forecast_high, agents_required_ceiling, pct_of_screening_staff)

```

```

tsa_dual <- tsa_agents_forecast %>%
  transmute(
    Month,
    `Passengers (Forecast + 10% Safety Buffer)` = pax_forecast_high,
    `Agents Required (rounded)` = agents_required_ceiling,
    `% of Screening Workforce Needed` = pct_of_screening_staff
  ) %>%
  pivot_longer(
    cols = -Month,
    names_to = "metric",
    values_to = "value"
  ) %>%
  drop_na(value)

ggplot(tsa_dual, aes(x = Month, y = value)) +
  geom_line() +
  geom_point() +
  facet_wrap(~ metric, scales = "free_y", ncol = 1) +
  labs(
    title = "2026 Resource Outlook: Forecast & Workforce Demand",
    subtitle = "Model: ARIMA(0,1,1)(0,1,0) | Metric: Forecast + 10% Safety Buffer",
    x = "Month",
    y = NULL
  ) +
  scale_y_continuous(labels = scales::label_number(accuracy = 0.01)) +  # <-- safer
  theme_minimal() +
  theme(
    strip.text = element_text(size = 12, face = "bold"),
    plot.title = element_text(face = "bold")
  )

```

2026 Resource Outlook: Forecast & Workforce Demand

Model: ARIMA(0,1,1)(0,1,0) | Metric: Forecast + 10% Safety Buffer

